

Survey on Conversion Measurement APIs

Joel Pfeiffer

Microsoft Advertising

Nomenclature

- **Features** – Contextual information about the click/impression
 - Keys
 - Sparse – CampaignId in space of uint64 values
 - Dense – Fixed dimension of continuous counts
 - $X := [\text{\#CampaignClicks}, \text{\#CampaignConversions}, \text{\#QueryClicks}, \text{\#QueryConversions}]$ (4Dim)
- **Labels** – Value assigned to a conversion
 - Boolean (conversion happened or not)
 - Counts of conversions
 - Conversion values (spend)
- **Trusted Helper**: Requires complete trust, will not reveal data and will not corrupt outputs
- **Semi-Trusted Helper**: Requires trust that the helper will not corrupt outputs, but does not require they will violate privacy (true data not revealed)
- **Malicious Helper**: Robust against a helper actively corrupting the outputs of the function

Order

- Browser Event Reporting APIs
 - Google Event Conversion API
 - Apple Private Conversion Measurement
- Browser Aggregate Reporting
 - Trusted third party server
 - MPC
 - Prio
 - Bucketization
 - Incremental DPF
 - Masked LARk
- Private Join Aggregate Reporting API (Meta / Mozilla)

Browser Event Reporting API (Generic)

- Browser level mechanism to attribute conversion to click
 - Extensions for impressions / views
- Core additional metadata attached to links
 - attributiondestination: eTLD+1 of the advertiser
 - attributionreportto: Usually, the ad network
 - attributioneventid: 64-bit event id
- Metadata stored in browser for click

... <user browses>

- User action on `http://<attributiondestination>/site`
 - Conversion pixel on page: `http://<attributionreportto>/conversiontracker`
 - Site redirects to: `https://<attributionreportto>/.well-known/attribution-reporting/trigger-attribution?trigger_data=trigger_data`
- Browser uses the `attributiondestination` and `attributionreportto` fields, determines the conversion happened for `attributioneventid`
 - Reports back to `attributionreportto` downstream what happened

Browser Event Reporting API

- Number of source event id bits varies between proposals
 - Can possibly be used for training
 - Generally, can be used for aggregation
- Number of trigger event bits varies between proposals
- Various random noise inserted
 - E.g., 5% random bits
- “Fire-and-forget”
- Timing delays / attacks
- For some of the proposals, (e.g., FLEDGE), some features not available

Order

- Browser Event Reporting APIs
 - Google Event Conversion API
 - Apple Private Conversion Measurement
- Browser Aggregate Reporting
 - Trusted third party server
 - MPC
 - Prio
 - Bucketization
 - Incremental DPF
 - Masked LARk
- Private Join Aggregate Reporting API (Meta / Mozilla)

Aggregate Reporting Helper Responsibilities

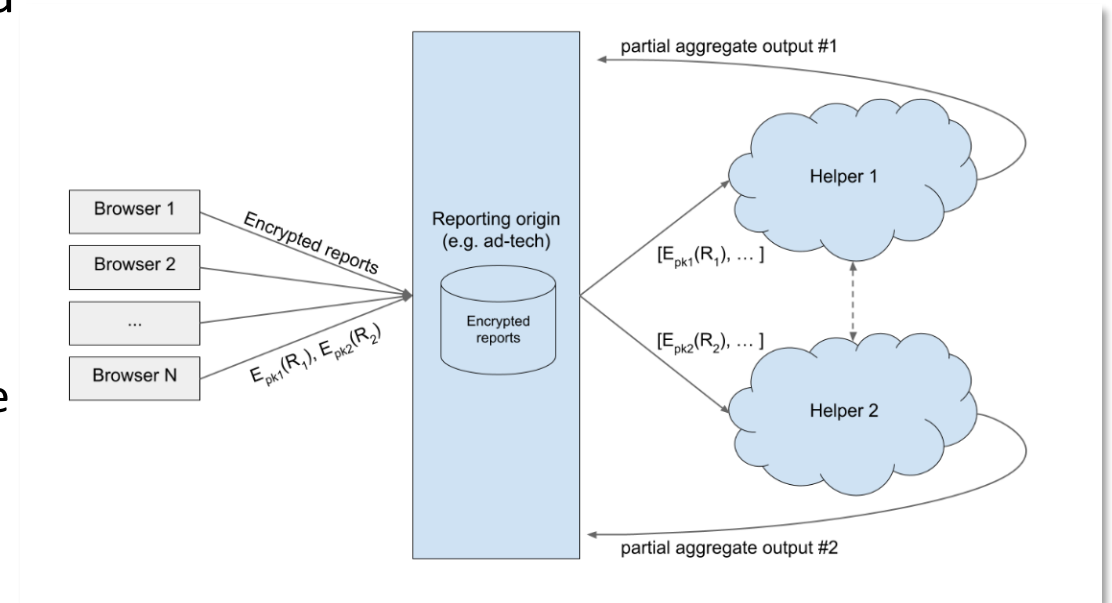
- Given encrypted data d and a function f , computes $f(d)$
- Responsible for **enforcing global privacy**
 - Differentially private outputs
 - K-anonymity
 - Privacy budget
- Most proposals are setup such that
 - Individual helpers cannot reveal user information, as they're always given a partial view of a user interaction
 - Helpers are trusted to return the right result, and not maliciously corrupting the outputs
 - This latter assumption can be periodically validated by advertising services

Trusted 3P Server

- Simply have a trusted third party, which agrees to not release information
- Trust verified through (e.g.) code audits, secure enclaves
- Queries are issued to the trusted server, which carries out the query
- Trusted server in charge restricting allowable queries, adding DP noise, etc.
- Extremely simple concept, easy to implement, efficient, etc.
- Single point of failure
 - Trust broken?
 - Hack / malicious outsiders have info in one place
- Cost / Supply

MPC Third Party Helpers

- Uses [Secure MultiParty Compute](#) to implement a trusted third party abstraction
- Rather than returning conversion information directly, returns a secret shared representation
 - Allow side information / contextual information
 - Half of the information is for [Helper 1](#)
 - Half of the information is for [Helper 2](#)
 - Looking at Helper 1 or Helper 2 alone appears to be random noise
 - Data for each helper is decrypted/encrypted with private/public keys
- Encrypted shares are given to the ad network
- Ad Network must call Helper 1 and Helper 2 to recover aggregates
- Helper 1 and Helper 2 enforce privacy

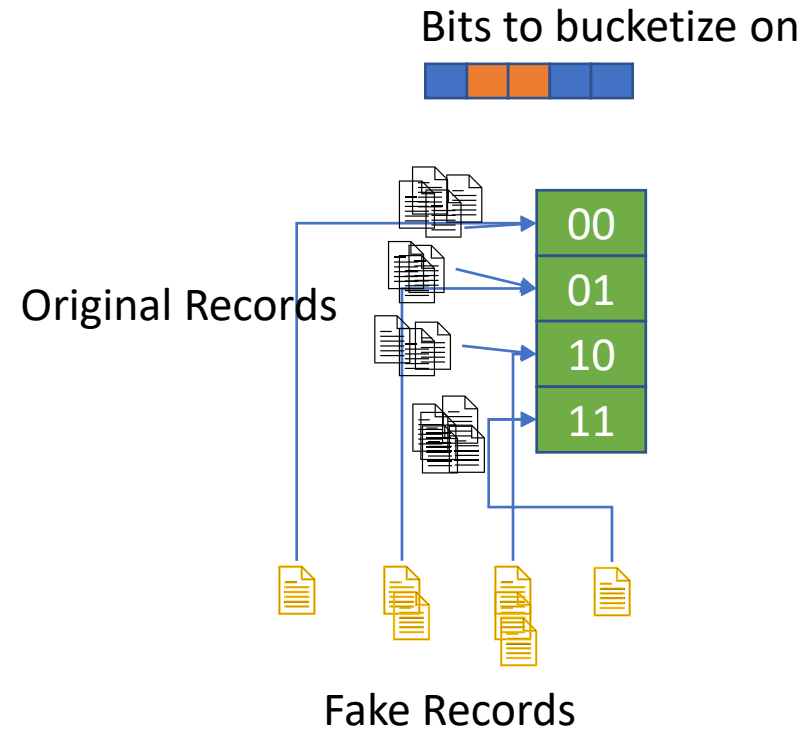


MPC - Prio

- Client secret shares a value intended for multiple servers
- Client secret shares side information to aggregate over
- In addition to secret shared side information, clients additionally provide Secret Shared non-interactive proof (SNIP)
 - Side information for servers to verify that they have a “Valid” sample, preventing poisoning attacks
- Extremely **efficient** for small number of bits
- Semi-trusted helpers
 - Features and Labels

Bucketization

- Handles aggregation and (in particular) aggregation on subsets
- Assumes some number of bits (e.g., 128)
 - Bits are Boolean shared between 2 parties
- High level idea
 - Take a subset of bits as the key
 - 2^{keybits} should be tractable
 - Each value of 2^{keybits} is called a “bucket”
 - Both helpers insert fake records proportional to 2^{keybits}
 - Shuffle and mask (3rd helper) to reassign keybits and obfuscate records from original 2 parties
 - Simply sum on the bits now to reconstruct
 - Can subsequently split on other bits




Bucketization

- Hides features / labels from Helpers
- Operates well within $2^{|keybit|}$ dense subsets
 - Small enough keybits should be tractable
 - Can operate on any subsets of bits
 - Can perform arbitrary hierarchical evaluations
 - Flexible ordering on the queries
- Needs (not required, but far, far more efficient) a third helper
 - Third helper can be chosen at evaluation time (not determined by browser)
- Semi-trusted helpers
 - Features and Labels
 - Some protection against malicious helpers

Incremental DPF (Google)

- Point Function
 - $x, y \in \{0, 1\}^*$ (e.g., 010011, 11011)
 - $P_{x,y}(x) = y$
 - $P_{\{x,y\}}(x') = 0^{|y|}$
- Distributed Point Function
 - Key pair k_0, k_1 s.t.
 - $F_{k_0} \oplus F_{k_1} := P$
- Incremental Distributed Point Function
 - Evaluate first b bits (using MPC) and get DP sum over all remaining bits
 - Can use b to help with $b+1$
 - Trim low occurrences to speedup



Incremental DPF

- Works well within sparse spaces
 - E.g., find 2^{16} most common out of 2^{128} space
- Fixed, hierarchical structure
- Semi-trusted helpers
 - Features and values

Masked LARk

- Intended as a private MapReduce framework
 - Browser is a map
 - Helpers are a reduce
- Primary example function is **training**, particularly SGD
- Browser creates multiple labels (**true** and **fake**) and sends both to the helpers (helpers are not told which are true and which are fake)
- Browser additionally provides secret shared **masks** which either include or exclude a sample from the computation
 - Masks appear random to the helpers
- Servers compute a function per sample, multiply each sample by its mask, and sum
 - Any bilinear function is fine, which includes summation
- Training requires a modification to always send some value, to handle the negative cases

Masked LARk

- Attacks
 - High cardinality label space – ad network could send ID as a label, then collude with helper
 - Randomly quantized labels
 - Requires default value sent if no conversion
 - Need to randomize / delay the record
- Non-communicative Helpers
 - Labels: Semi-trusted helpers
 - Features: Trusted Helpers
 - Models: Trusted Helpers

Private Joins: Meta Proposal

- Separates out *click* app versus *conversion* app
 - Focuses on LinkedIn/Facebook user view resulting in an Edge/Chrome conversion
- In particular, using a common API, apps and websites will be able to set and use a **match key**
 - This match key can be set per site, similar to a dictionary:
{ "linkedin.com" : "match_key_li" }
 - These can be used to share event information, but not readable outside a common API
 - Presumably user must be signed on in both locations to set the keys
- Data, either click or conversion, is encrypted using the API and corresponding match key for the reporting site, and sent to the advertising service

Private Joins: Meta Proposal

- Clicks and conversions are subsequently joined using MPC
- Join keys are blinded with homomorphic encryption to destroy the match keys being read by MPC
 - Data is then shuffled
 - Joining done with a garbled circuit and time ordering
- Aggregation can now be done
- Semi-trusted helpers

References

- Google Conversion Event API: <https://github.com/WICG/conversion-measurement-api/blob/main/EVENT.md>
- Google Aggregate API: <https://github.com/WICG/conversion-measurement-api/blob/main/AGGREGATE.md>
- PCM: <https://webkit.org/blog/11529/introducing-private-click-measurement-pcm/>
- Masked LARk
 - Explainer: <https://github.com/WICG/privacy-preserving-ads>
 - Paper: <https://arxiv.org/abs/2110.14794>
 - Prototype: <https://github.com/microsoft/MaskedLARk>
- Bucketization:
 - <https://eprint.iacr.org/2021/1490>
- IPA:
 - Prototype: <https://github.com/martinthomson/raw-ipa?ref=rustrepo.com>
 - Paper: <https://docs.google.com/document/d/1KpdSKD8-Rn0bWPTu4UtK54ks0yv2j22pA5SrAD9av4s/edit#heading=h.f4x9f0nqv28x>