

# Masked LARk: The Masked Learning, Aggregation and Reporting workflow

Joseph Pfeiffer, [joelpf@microsoft.com](mailto:joelpf@microsoft.com)

Denis Charles: [cdx@microsoft.com](mailto:cdx@microsoft.com)

Mehul Parsana\*: [mparsana@microsoft.com](mailto:mparsana@microsoft.com)

Erik Anderson: [erikan@microsoft.com](mailto:erikan@microsoft.com)

Young Hun Jung: [youjung@microsoft.com](mailto:youjung@microsoft.com)

Davis Gilton: [davisgilton@microsoft.com](mailto:davisgilton@microsoft.com)

---

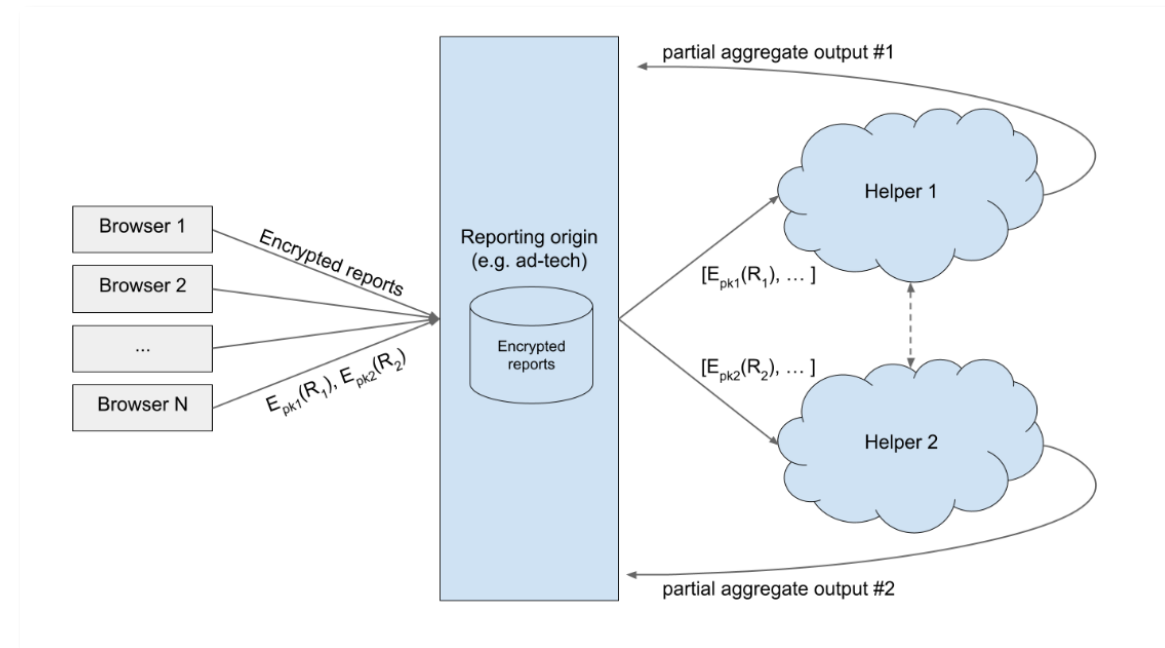
# Agenda

- Related Work:  
Google (Browser) Aggregate Conversion  
MultiParty Compute Workflow
- Masked LARk: Masked Learning  
Aggregation and Reporting workflow
  - Workflow Overview
  - Gradient Computation
- Experiments
- Conclusions

# Related Work:

## Google Proposal for Conversion Reporting

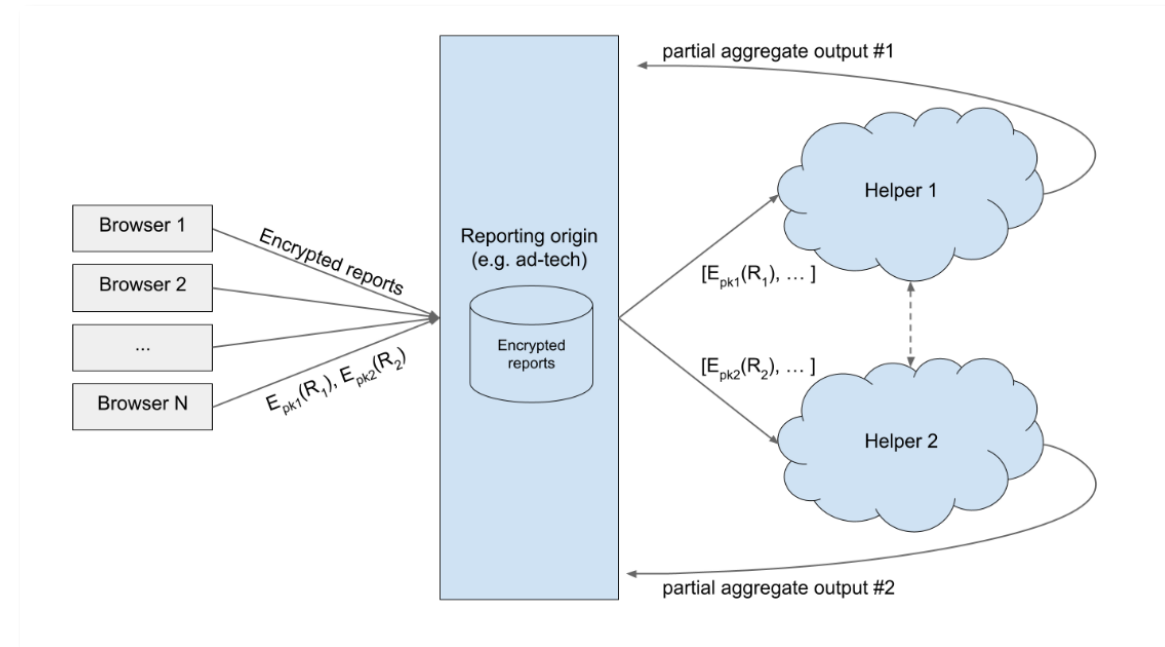
- Recent Google Proposal discusses utilizing secure MultiParty Compute (MPC) for Trusted Mediator\*
- MPC performs both aggregation and enforces differential privacy constraints
- Upon leaving browser, no **individual entity** has a complete picture of an individual record
  - Secret Sharing



# Related Work:

## Google Proposal for Conversion Reporting

- Pros:
  - Segregated helpers implementing aggregation more palatable than single trusted mediator
- Limitations:
  - Handles aggregation for some reporting needs but does not address modeling





# Masked LARk: Masked Learning Aggregation and Reporting workflow

# Masked LARk

- **Goal:** Expand the aggregation service towards an abstracted differentially private **Map-Reduce** framework
  - **Browsers** implement a secure “Map” function
  - **Helpers** implement a differentially private “Reduce” operation
  - **Advertising Servers** are the consumers and data storage
- **Goal:** Expand from single aggregation function, to a platform that enables a variety of differentially private functions (e.g., Aggregation, **Model Training**)
- Provided
  - Explainer: <https://github.com/WICG/privacy-preserving-ads/blob/main/MaskedLARK.md>
  - Paper: <https://arxiv.org/abs/2110.14794>
  - <sup>6</sup> • Prototype: <https://github.com/microsoft/MaskedLARK>

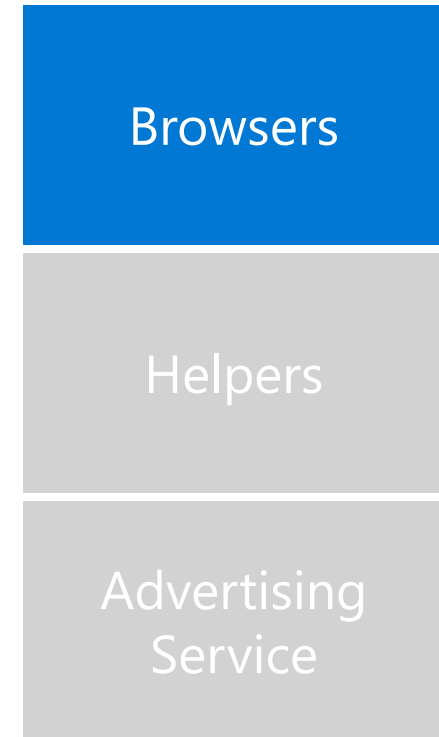
Browsers

Helpers

Advertising  
Service

# Masked LARk - Browsers

- Browsers represent the user interest
  - Initial viewing / clicking of advertisement
  - The conversion
- Only party with **full information** about individual users
- Responsibilities
  - **Attributing a conversion to a view or click**
  - Inserting fake records and masks for true or fake records
  - In certain scenarios, inserting local differential privacy
  - Secret sharing of values for helpers to aggregate / train on
  - Sending reports to the advertising server, which are encrypted so only the helper can read
  - Choosing which helpers to use





# Masked LARk - Helpers

- Receive encrypted reports from advertising servers
- Helpers perform core function implementation, doing most of the heavy lifting within the platform
  - E.g., aggregation or model training
- Enforce community-approved privacy constraints
  - *K-Anonymity*: require a certain minimum number of records to return any results
  - *Global differential privacy* added to outputs of functions
  - *Manage privacy budget*

Browsers

Helpers

Advertising  
Service



# Masked LARk – Advertising Server

- At impression time, passes relevant features to Browser for later processing
- Acts as a data storage unit
  - Records received from Browsers are encrypted with keys held by the helpers
  - Advertising Service holds these for later processing, bearing this cost
- Can only retrieve aggregation information utilizing the associated helpers
- Applies aggregated information to future tasks
  - Reporting conversion data to advertisers
  - Model training and later inference

Browsers

Helpers

Advertising  
Service



# Implemented Helper Functions: Gradient Computation

# Model Training

- Want to allow Masked LARk to handle complex scenarios, specifically model training, where the features are observed yet the labels are secrets
- Core issue:
  - Most differentiable models (e.g., neural networks) optimize some variation of SGD

$$\theta_{j+1} = \theta_j - \eta \sum_i \nabla \text{Loss}(\text{features} = x_i, \text{label} = y_i, \text{Model} = \theta_j)$$

- $\nabla L(x_i, y_i, \theta_j)$  is computed per sample, requiring the features and labels together
  - As is, this would reveal features and labels to the helper
- Will show, can utilize MPC with masking to compute the gradient step

# Model Training – Key Idea

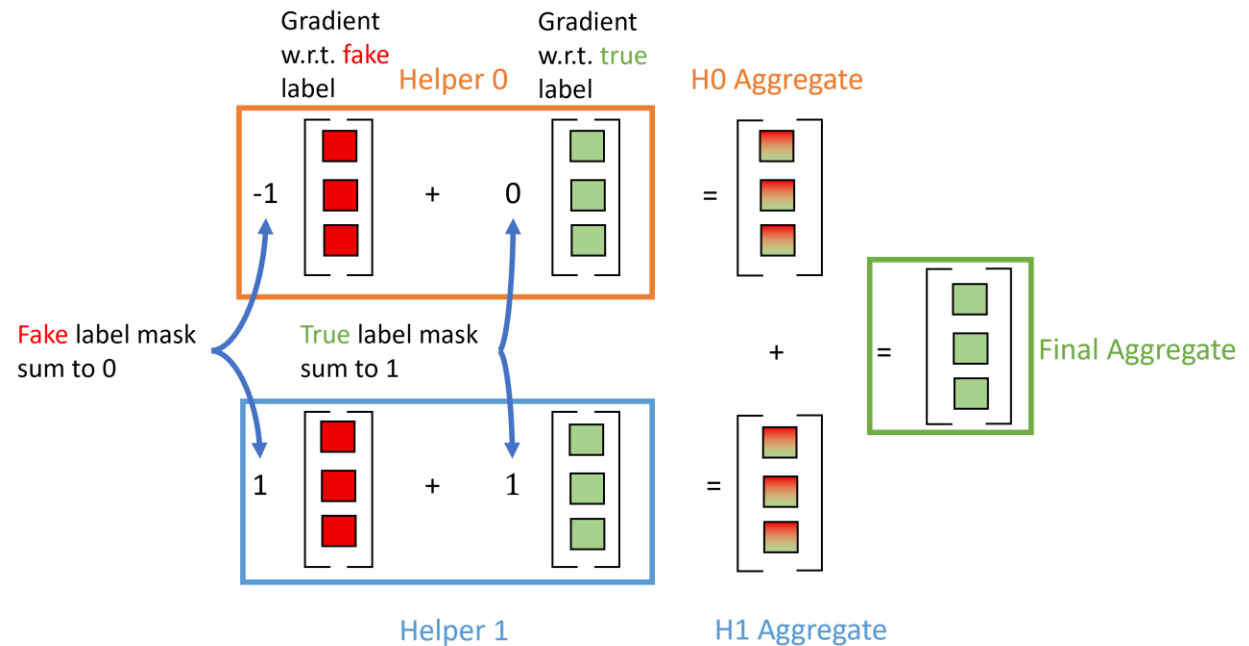
- Browsers will generate many possible labels / features, some which may be true but will contain some that are “fake” (i.e., not the correct label / feature)
- Both helpers will compute gradients for both the true / fake examples
- Browser will additionally send masks  $\alpha$  to the helpers to include in the summation

$$\sum_i \alpha_i \nabla \text{Loss}(x_i, y_i, \theta_j)$$

- $\alpha$  are constructed in such a way that fake examples are removed from the final MPC summation

# Example: Single Sample Mask (Proof Sketch)

- Assume
  - One Sample, Two Helpers
  - One **True** Label (e.g., conversion occurred) and One **Fake** Label (e.g., no conversion occurred)
- Each helper computes gradient for both **True** and **Fake** Labels
- Each Helper multiplies each gradient by the matching mask
- Each Helper sums the vectors
- Advertising server sums the partials returned by the Helpers



# More general formulation

- Let  $X \subset \mathcal{X}$ ,  $Y \subset \mathcal{Y}$ ,  $\theta \in \Theta$  indicate features, labels and a model
- Let  $f: \mathcal{X} \times \mathcal{Y} \times \Theta \rightarrow \mathbb{R}^d$  be a mapping from features, labels and models to some vector space
  - Focus on gradient computation for  $f$ , but other functions could potentially apply
- Let  $g$  indicate a bi-linear aggregation function over a set of samples after applying  $f$ , e.g.
  - $g(\langle f(x_1, y_1, \theta), \dots, f(x_n, y_n, \theta) \rangle) := \sum_{i \in n} f(x_i, y_i, \theta)$
  - $g$  satisfies bi-linearity, i.e., for  $\alpha_i, \beta_i \in \mathbb{R}$ :
    - $g(\langle (\alpha_i + \beta_i) f(x_i, y_i, \theta) \rangle) = g(\langle \alpha_i f(x_i, y_i, \theta), \beta_i f(x_i, y_i, \theta) \rangle)$
  - Focus on summation for  $g$ , but other functions could potentially apply

# Summation over Masked Gradients equals summation over true gradients

- Section 4.2, Corollary 4.1.1
  - For each helper  $h$ 
    - $\Psi^h := \sum_i \alpha_i^h \nabla \text{Loss}(x_i, y_i, \theta_j)$
    - $\sum_i \nabla \text{Loss}(x_i, y_i, \theta_j) = \sum_h \Psi^h$
- Generalizes to arbitrary mappings  $f$  and all bilinear functions  $g$



# Fake Labels and Masks

- Let  $\bar{y}_i \in \mathcal{Y} \setminus \{y_i\}$  be a fake label additionally sent from the browser
  - Cannot be  $y_i$
  - Both fake and real can be sent from the browser, but the order is randomized
    - The corresponding masks (to be defined) must match the randomized order
  - For real-valued labels, the values are quantized
  - For non-integer values, the values are randomly rounded
- Without loss of generality, define masks for the real/fake labels for two helpers:
  - $\alpha_i^0 + \alpha_i^1 = 1$  [True label]
  - $\bar{\alpha}_i^0 + \bar{\alpha}_i^1 = 0$  [Fake label]
  - $\alpha_i^0, \bar{\alpha}_i^0$  are both random variables on a finite ring, solve for  $\alpha_i^1, \bar{\alpha}_i^1$  as appropriate

# Threat Models - Attacks

- High cardinality label space (e.g., floats) – ad network could send ID as label (lower order bits, etc), then collude with helper to recover features and label
  - Randomly Quantized labels to a desired cardinality
  - In expectation, minimizes same loss
- Poisoning by the Browser (sending invalid label masks)
  - Some overhead for a validity function (e.g., SNIPs)
- Requirement of default (0) label when expired
  - Randomized ending times
- One Helper colludes w/ advertising service
  - Features exposed: In some situations, features could be considered sensitive
  - Add random noise to feature vector (local DP)
  - Extra cost: Can secret share  $x = x_0 + x_1$ , as  $W \cdot x = W \cdot x_0 + W \cdot x_1$

# Efficiency-Privacy Tradeoffs

Efficiency/Utility

Privacy



## Event Conversion API

Plausible Deniability  
Limited Information

## Masked Event Conversion API

Rather than limit bits, return multiple possible (quantized) labels.  
Keys / features derived only from first party information  
Also provide secret shared masks for both true and false labels.  
Requires minimal MPC, can do various partitioning / slicing / any model training as needed  
Heavy computation (gradient) can be performed on ad side

Obvious timing attacks

## Masked LARk

Advertising Server cannot recover features w/o collusion  
Helpers cannot recover labels w/o collusion with each other  
Fairly efficient (communication)  
Insert local DP

## Buckets/iDPF

Secret Share Feature Vector  
Aggregation scope

## SS Masked LARk

Secret Share Feature Vector  
Much more compute  
Does not necessarily extend to all bilinear functions



# Experiments

# Additional Privacy Requirements and Implementation Details

- Tested Local Differential Privacy for the feature vector
- Malicious advertising server and helper could reverse engineer label from a gradient and a feature vector (4.2.2)
  - $k$ -anonymity
  - Global differential privacy (Gradient Clipping & Laplace noise)
- Feature vector assumed small and dense (e.g., 100 dimensional vector)
  - Feature vector assumed bytes to save space

Code: <https://github.com/microsoft/MaskedLARk>

# Implementation: PyTorch-friendly Training Interface

```
import MaskedLark as mlark
```

```
# Acts as an interface to Azure-based helper services
```


```
helper = mlark.Helper(helper_config)
```

```
for ii, sample_batch in enumerate(dataloader):
```

```
    helper.fetch_gradients(model)
```

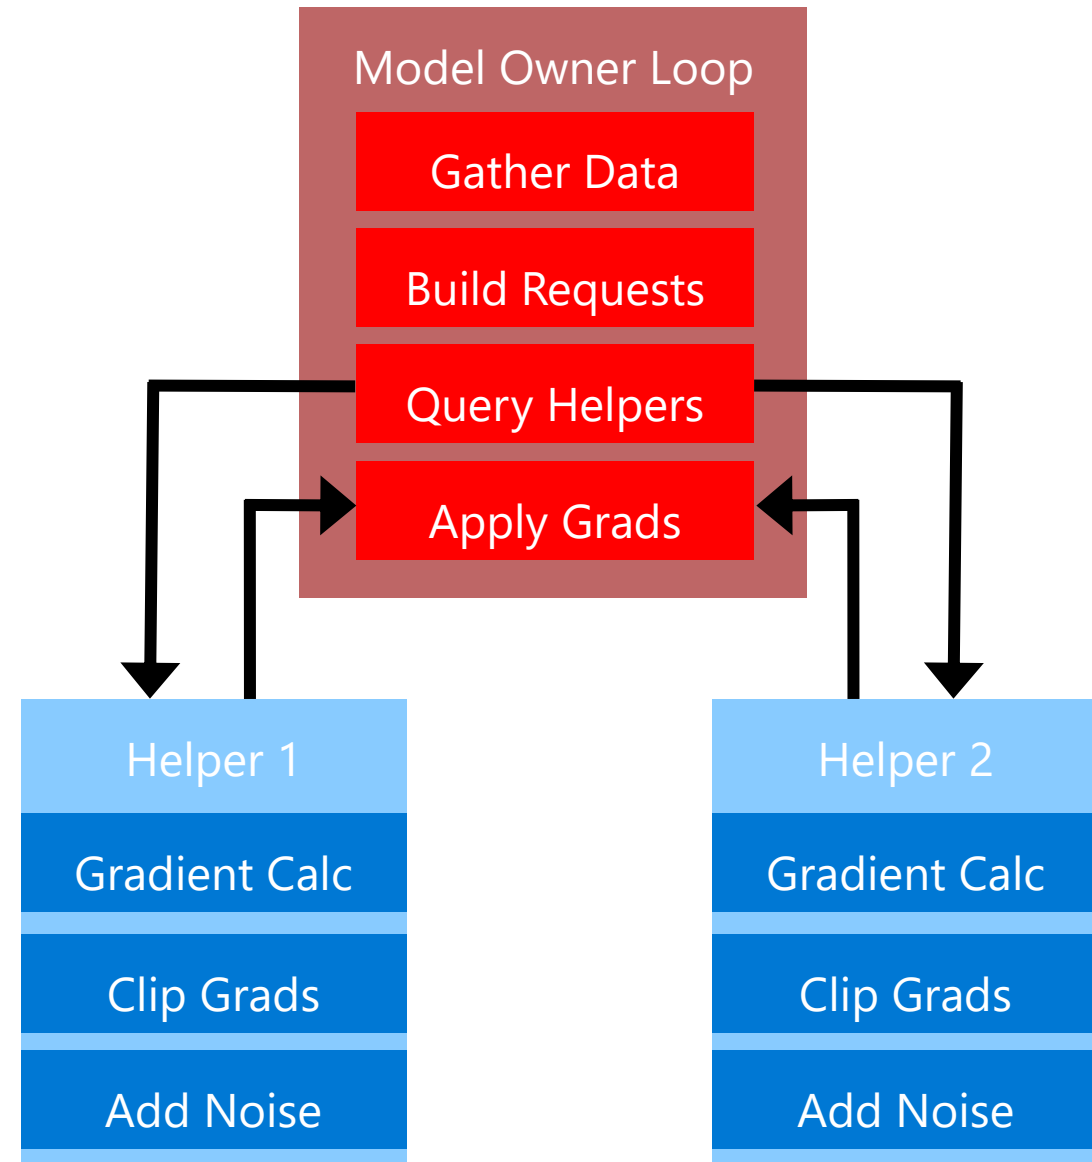
```
    helper.backward()
```

Code: <https://github.com/microsoft/MaskedLARK>



# Experiment Setup


- Test scenario: model training
  - Helpers are simulated with Azure Functions and called via our API
  - Helpers use PyTorch's autograd to calculate gradients
- Models are passed as serialized ONNX models.
- We test on MNIST (784 dimensions) and WBCD (30 dimensions)
  - MNIST model: one hidden layer with 500 nodes.
  - WBCD model: two hidden layers with 50 nodes.






# Lesson: Privacy Mechanisms Affect Model Performance

Aggressive gradient clipping hurts performance, but networks are resilient to small amounts of noise.




MNIST




WBCD

# Lesson: When Training Remotely, Use Large Batch Sizes

Training using helper services incentivizes sending large amounts of data at one time.




MNIST




WBCD

# Alternate Option: Local Differential Privacy

Gradient clipping can be costly: we can try adding DP noise to the inputs. Networks are surprisingly resilient to this type of noise!



MNIST



WBCD



# Conclusions

# Conclusions

- Designed platform, Masked LARk, for secure MPC and differentially private aggregation and model training
- Outlined the core functionality and algorithms, along with extensions for additional user or model privacy
- Implemented helper services within Azure, notably the gradient computation for model training
- Analysis on two publicly available datasets
- Released at <https://github.com/microsoft/maskedlark>
- arXiv version of paper - [link](#).



# Q&A