



Minimal Viable Functionality for Measurement

Mariana Raykova
Google



Why common API?

A common API that is supported across different browsers and apps is an important goal that will most effectively meet the web ecosystem needs

- There seems to be agreement



Why common API even if restricted?

An API with the above properties is important even if it **does not meet everybody's complete set of requirements** (*different parties may have different extensions simultaneously*) :

Pros:

- Implementing and deploying something simpler will surface system challenges, which can be resolved before building complexity
- This is a first step, which can be extended later

Cons

- Modifications to the V1 may be hard and take a long time (we should think how it will be easy to iterate)
- Different parties may have different notions of MVP - we need agreement



Functionality Requirements

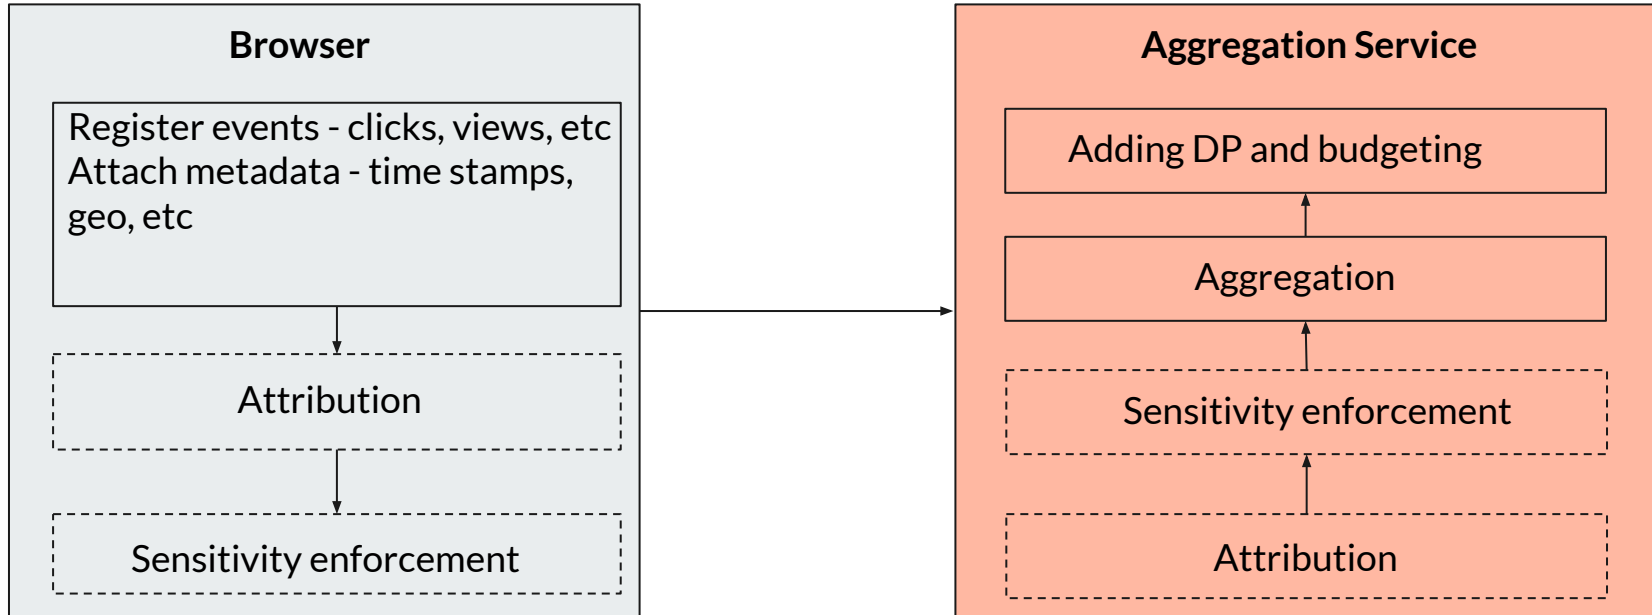
Aggregation vs Machine Learning

- Aggregation is useful on its own - count and value, building block for some simplifies ML
- Machine learning functionality will be needed to address all use cases eventually

Attribution on Device or off Device

- On device reduces complexity and security perimeter that needs to be protected with privacy preserving computation
- Off device offers cross-device attribution which is valuable in many settings

Strawman proposals - aggregation only





Most basic strawman functionality

Functionality:

- Only aggregation with DP
 - fixed number of aggregation buckets on the order of 2^{20} - 2^{32}
 - Sensitivity bound is fixed and applied on device
 - Adaptive epsilon per query
- Attribution on device


In the context of existing API proposals - differences:

- Chrome Attribution API - many more aggregation buckets 2^{128}
- Apple PCM - only 8 buckets and no DP
- Mozilla-Meta IPA - attribution off device on the servers
- Microsoft measurement proposal - more flexible aggregation buckets



Missing functionality

- Adaptively defined aggregation buckets at time of query
 - Importance - the aggregation granularity can be informed by different queries, e.g. for campaigns with many counts, consider finer-grained aggregation with attributes
- Different sensitivity bound per query
 - Different sensitivity bounds are optimal for different functionalities
 - Criteo challenge - logistic regression training: L2 bound $\sqrt{190}$, L1 bound 190
 - Count and value aggregates - L1 bound better result if the mass is on one key



Strawman functionality - adaptive query and DP

Functionality:

- Only aggregation with adaptive queries and dynamic DP budget per query
 - A query is specified by a set of **bucketing functions** F_1, \dots, F_n (chosen together), a **sensitivity bound** (e.g. L_0, L_1, L_{inf}) for the record contributions and an **epsilon value**
 - The aggregation uses the following buckets: each record is mapped to a bucket $F(\text{record impression data})$
 - The aggregators enforce the sensitivity bound for each record during aggregation
 - DP noise is added to the aggregated buckets according to the specified epsilon
- Attribution on device



Optionality for Extensions

- “High-level architecture of the API” allows flexibility for future use cases - how do we formalize this requirement?
 - Simple ML like logistic regression supported through the flexible aggregation
 - If we have an architecture where we are sending secret shared data, changes in functionality affect mostly the aggregation service side.
 - We should aim to minimize what goes out of the browser, but this can be changed in this architecture



Threat model

- Initial Proposal
 - Two parties - minimizes trust assumption
 - Malicious for privacy but semi-honest for correctness
 - Relaxation allow limited DP leakage to servers in the form of DP counts associated that cannot be linked to meaningful other impression, conversion data
 - Client input authentication - optional, can have approximation with trust tokens
- Cost - come up with a measure of acceptable cost